


Why the accuracy curves appear to live above 0.4



The y-axis in every accuracy panel of the study is $[0, 1]$, not $[0.4, 1]$ -- accuracy is a proportion and the full axis is drawn deliberately, so that differences of a tenth of a percentage point are not visually inflated; each panel carries a zoom inset for the real resolution. (A) Accuracy of a random initial state on TITANIC over 200,000 draws uniform in the unit ball: mean 0.4998, SD 0.1127, minimum 0.2022. The shaded strip is mean $\pm$ 1 SD, whose lower edge is 0.387 -- that is the apparent 0.4. 25% of draws fall below 0.40 and no figure in the study draws them. (B) The same reversible run whose band appears in the main figures, with the full min-max envelope added: the lowest state any replicate visits is 0.274, far below the band. (C) Both constraint sets are symmetric in w , and $\text{acc}(w) + \text{acc}(-w) = 1$ EXACTLY -- checked here to 0e+00 over 200,000 draws on each dataset, because a linear rule and its negation disagree on every point with $x \cdot w \neq 0$. So each final state has an equally feasible mirror at accuracy 0.211. The geometry does not rule the low-accuracy half out; the potential does, by 407 units of U . The mean at iteration 0 is exactly 1/2 for the same reason: the initial law is symmetric under $w \rightarrow -w$, so accuracy pairs up around 1/2 whatever the data look like.