



titanic_lp, $\eta = 3 \times 10^{-5}$, $s = (2, 2, 2)$: the post-burn-in estimator



R = 1500 coupled replicates, seed 5200 (disjoint from the screen above), first half of each 2000-step run discarded as burn-in (declared in advance). Each replicate contributes one number per panel: its accuracy averaged over the checkpointed post-burn-in states. The two arms share W_0 and the Gaussian stream, so the difference is paired and the histograms overlap far more than the paired s.e. suggests -- the spread here is across replicates, the s.e. is across PAIRS. The vertical lines are the arm means. This is the same stationary quantity the final-iterate estimator targets, computed with lower variance; both arms provably share that target, so a difference here is a difference in estimator variance and in finite-eta bias, not in the target.