


The efficiency gain is a short-budget gain, and it does reverse



RMSE of the estimated posterior mean, pooled over 12 test functions, against compute budget. At each budget BOTH arms are re-optimised over their own (η , s). The non-reversible arm is ahead at short budgets ($\times 1.33$ lower RMSE at $T = 1$) and the curves cross near $T = 10$, after which the reversible arm can afford a small enough step that its lower bias wins. This is the honest shape of the result and the single most important caveat on it.