



titanic_lp_anchored_accuracy



TITANIC -- smoothed ℓ_p $g(\beta) = \sum_i (\beta_i^2 + \epsilon^2)^{p/2} \leq \Lambda$, $p = 2.4$, $\epsilon = 0.18$, $\Lambda = 4$. $d = 9$ coefficients, no intercept. Step size $h = 0.0001$ (actually used). Mini-batch $m = 30$. $R = 100$ independent sampler replicates on one fixed stratified 80/20 split ($n_{\text{train}} = 712$, $n_{\text{test}} = 179$). Anchor parameter $\rho = \log 2 = 0.693147$ for the anchored methods ($\rho = 0$ otherwise). Block strengths $s = (1, 0.25, 1)$. Shaded bands: mean ± 1 SD (ddof = 1) across replicates -- repeat-run variability conditional on the fixed dataset and split, not a confidence or credible interval. Block strengths selected by a training-only sweep over 14 candidates (criterion: mean TRAINING accuracy over the last 25% of checkpoints); test labels were not used to choose s .