



titanic_lp: anchored Langevin, s = 5



TITANIC -- smoothed ℓ_p $g(\beta) = \sum_i (\beta_i^2 + \epsilon^2)^{p/2} \leq \Lambda$, $p = 2.4$, $\epsilon = 0.18$, $\Lambda = 4$. $d = 9$, no intercept. Block strengths $s = (5, 5, 5)$. Step size $h = 1e-05$ (actually used). Mini-batch $m = 30$. $R = 100$ independent replicates on one fixed stratified 80/20 split ($n_{\text{train}} = 712$, $n_{\text{test}} = 179$). Anchor $\rho = \log 2 = 0.693147$ for both arms. Bands: mean $\pm$ 1 SD (ddof = 1) across replicates -- repeat-run variability at a fixed split, not a confidence or credible interval. h selected per experiment by the paired TRAINING gain over 3 candidates; these curves are a CONFIRMATION run on independent replicates (seed 4100) not used for that selection. Verdict on the paired training difference: non-reversible WINS.