



Non-reversible vs reversible anchored Langevin — L1-smooth ball

Training accuracy (n = 1600)



Test accuracy (n = 400)



d = 9; constraint: L1-smooth ball, $g(\beta) = \sum_i \sqrt{\beta_i^2 + \epsilon^2} \leq 3.70$ ($\epsilon = 0.2$, $l1_radius = 1.9$); AR(1) predictors $\rho_x = 0.99$ (posterior condition number 1225);
step size $h = 0.0001$; mini-batch $m = 50$; repeats $R = 100$; iterations = 1000; anchor $\rho = \log 2$ for BOTH methods; block strengths $s = (2.0, 2.0, 2.0)$.

Lines are the across-replicate mean; bands are mean ± 1 sample sd (ddof = 1) of single-iterate accuracy, i.e. repeat-run variability -- NOT confidence or posterior credible intervals. The vertical axis is windowed on the plateau, not $[0, 1]$.