



Titanic — non-reversible vs reversible anchored Langevin (l1 constraint)



d = 9 (intercept + 8 columns: emb_C, emb_Q, sex, pclass, age, sibsp, log_fare, parch); n_train = 712, n_test = 179; $U = f + g$, $g = 2 \cdot \sum |w_j|$, anchor delta = 0.02; EXACT gradient; eta = 2.0e-05; s = 4; R = 100. Preprocessing: scaling = 'design', order = 'design', signs aligned = True, block reparametrisation (v_axis, v_fast, v_slow) = (1, 64, 2), (b, e) = (1.0, 0.5). Reference (scikit-learn logistic regression) test accuracy 0.782. Lines: across-replicate mean; bands: mean +/- 1 sd of single-iterate accuracy (repeat-run variability). Shared initialisation and noise per replicate. y-axis fixed to (0.4, 0.9), not windowed.